

Week 1 Setting up a data science working environment

- Key tools:
 - the anaconda environment
 - vscode
 - Rstudio
- Projects
 - projects, files, directories
 - the terminal
 - self documenting code
 - * jupyter notebooks, markdown, and R markdown
 - essentials of REPL python in jupyter
 - * numeric variables, arithmetic, scientific functions;
 - * lists and arrays; indexing and slicing; vectorization
 - basics of R
 - * datatypes, arrays, arithmetic, slicing and indexing, vectorization
 - a minimal introduction to plotting in R and Python

Week 2 Probability and Statistics: a first look at the normal distribution

- Working with normally distributed populations
 - Events and Outcomes
 - The Normal Distribution
 - * probability mass functions and area under the curve
 - Mean and Variance
 - Hypothesis testing
 - * sampling from a univariate normal distribution
 - * Null and alternative hypotheses
 - * p-values, statistical significance and confidence intervals (for univariate normally distributed populations)
 - illustrated with examples in R and Python

Week 3 Programming with Data in R and Python

- Key tools
 - R dataframes and the tidyverse
 - Numpy and Pandas
 - functions in R and Python
- working with files and I/O in R and Python
- pandas Series and dataframe basics (reading files, indices, selecting, summarizing data)
- R factors and dataframes (reading files, indices, selecting, summarizing)

Week 4 Discrete Probability and Bayes Theorem

- Discrete probability;
 - events and outcomes;
 - mean and variance
 - independent events;
 - conditional probability and Bayes theorem in the discrete case
 - bernoulli and binomial distributions
 - false positives, false negatives, versions of the base rate fallacy
 - discussion of the Naive Bayes spam filter (?)
 - illustrated with R and Python examples

Week 5 Slicing and dicing data in R and Python

- One day on pandas grouping, summarizing, selecting data
- One day on R grouping, summarizing, selecting data (tidyverse)

Week 6 Linear Algebra

- Geometry of n-dimensional space, vectors, addition and scalar multiplication of vectors, the dot product, orthogonality
- Matrices, matrix multiplication, column space of a matrix
- Ordinary Least Squares as an illustration(?) of geometry and linear algebra
- computational examples in both R and Python (numpy)

Week 7-8 Multivariate calculus

- discussion of functions of several variables:
 - graphs of functions
 - contour graphs and level surfaces
- review of the derivative in one dimension; rates of change
- partial derivatives
- directional derivatives and the gradient
- relationship of the gradient to the level curves
- Use of calculus to solve the OLS problem done by linear algebra in week 5 (exact solution)
- Use of gradient descent to solve OLS problem
- includes programming examples for gradient descent in the least squares case

Week 9 A deeper dive into Visualization

- more on ggplot and its capabilities
- python plotting packages (seaborn? bokeh?)

Week 10 Statistical models

- What is a statistical model?
- Likelihood and model parameters
- Maximum likelihood estimation and gradient descent
- Illustrated with OLS and Logistic Regression?

Week 11-12 More on Statistics

- Another look at the normal distribution; the multivariate normal
- covariance, correlation
- estimation of parameters for linear regression and logistic regression
- significance and confidence intervals
- null and alternative hypotheses
- p-values

Week 13 Version Control

- Git as a tool (command line and through R studio)
- commits, branches
- remotes and github
- using github to host a web page for a project
- collaboration using Git – pull requests; contributing to open source projects

Week 14 Advanced topics in programming

- Data structures; object oriented concepts
- Essential notions from data structures:
 - stacks, lists, hashing
- Python classes
 - data attributes and methods

Week 14 Databases

- What is a relational database? tables, keys, indices, joins
- Basic SQL for getting data from a database